# CW2 Plan

## 1 Goal

**CW2 output:** 10-minute Dragons' Den pitch + 5-minute Q&A, plus a **GitHub evidence pack** (slides/videos/prototype/testing).

**70+ recipe:**

- **Execution & tangible outputs (30%):** functional MVP / pilot, clear progress beyond CW1.
- **Testing, evidence & iteration (20%):** real user testing, measured learnings, iteration decisions, GenAI disclosure.
- **Semester 2 concepts (15%):** legal + ethical + sustainability + international considerations *affect decisions* (not just mentioned).
- **Reflection & forward strategy (15%):** honest constraints/risks + realistic next steps.
- **Communication (15%):** polished, multimodal pitch + strong Q&A handling.
- **Group contribution (5%):** clear evidence of equitable teamwork.

## 2 What we will build (MVP v1 that looks impressive and is demo-able)

### Principle: ship a tight end-to-end workflow

Do **one core workflow** extremely well; avoid spreading effort across too many features.

### 2.1 MVP v1 (minimum that still scores highly)

**Core flow (must work end-to-end):**

- Project creation (*Create project*)
- Add a paper: upload PDF *and/or* add by DOI (even a simple stub is OK)
- Split view: PDF reader + notes
- Highlight text → *Add to notebook* (with tags)
- Notebook sections (e.g., Intro / Lit Review / Methods) linked to sources
- **Scoped AI:** answers only from the project corpus **with citations back to source passages**
- **Dynamic bibliography** per notebook section (basic is fine if consistent)
- Nice-to-have (only if time): Simple knowledge graph view that links concepts → citations/sources

### 2.2 Legal/IP constraint (do this to score higher)

If CW1 implied auto-unlocking paywalled PDFs via institutional SSO, for CW2 we should **explicitly pivot**:

- **No auto-download of paywalled PDFs.**
- Use **link-out** to publisher/university access routes OR rely on **user-uploaded PDFs**.

This becomes a strong "Semester 2 decision" slide: *we identified legal risk and changed product design accordingly.*
*...the above can be changed if we find something else better*

## 3 Testing & iteration (how to make it look like real entrepreneurship)

We will run **two rounds of testing** and show concrete iteration.

### 3.1 Round A: quick usability test (n=5–8)

**Tasks (give each tester 3 tasks):**

1. Import paper (upload/DOI), open reader
2. Highlight a claim and add to notebook with a tag
3. Ask AI for a short paragraph *with citations* and paste into notes

**Metrics to capture:**

- Time-to-complete per task
- Where they got stuck (notes)
- 1–5 usefulness score + 1–5 trust score

**Iteration:** fix top 3 issues. Document: *feedback* $\rightarrow$ *decision* $\rightarrow$ *product change*.

### 3.2 Round B: mini-pilot (n=3–5 for 2–5 days)

**Pilot prompt:** use ScholarStack for real coursework reading for a few days.
**Evidence:**

- Daily short log (2 min/day) + final interview
- Final satisfaction + "would you continue using?" (Y/N + why)

### 3.3 Evidence artifacts (make markers see receipts fast)

- Consent script + ethics note (simple, clear)
- Anonymised notes + screenshots
- 20–60s screen recordings of users completing tasks (with permission)
- Before/after change log

## 4 Market / launch experiments (quick traction numbers)

Do **one small experiment that yields real numbers**:

- Landing page + waitlist (A/B two versions of headline/value prop)
- Share to relevant groups (course cohorts, dissertation groups, PhD/postgrad channels)
- Track: visits → signups → demo requests (conversion rate)

*Could use a site to develop a quick landing page.*

## 5 Semester 2 concepts (easy way to score the 15% without waffle)

Pick **3–4 issues** and show the **decision** each forced.

- **Legal/IP:** respect publisher licensing (no auto-download); link-out + user upload.
- **Data protection (GDPR):** retention limits, deletion, encryption; clarify what is sent to LLM provider.
- **Ethics/trust:** hallucination risk → scoped RAG + citations + UI warnings.
- **Sustainability/cost:** token limits, caching summaries, "AI only when needed", tier gating.
- **International:** different uni auth/licensing → modular auth + localisation plan.

  *...the above are suggested in the spec, fill out more as the semester progresses*

## 6 Pitch structure (10 minutes)

**Slide order:**

1. Problem + why now
2. What we built (MVP) + a 1-line value proposition
3. Live demo teaser → **60–90s end-to-end workflow demo**
4. Testing evidence: Round A → iteration → Round B pilot results (numbers + quotes)
5. Market/traction: waitlist + who converted best
6. Business model/pricing: tiers + willingness-to-pay learning
7. Constraints that shaped product: legal/ethical/sustainability (one slide)
8. Risks + mitigation (short, credible)
9. Next 8 weeks roadmap (realistic)
10. Closing + Q&A prompt

## 7 Q&A prep (questions we must be ready for)

- "How are you different from existing tools (NotebookLM / Zotero / Readwise etc.)?"
- "How do you handle paywalled papers legally?"
- "AI costs per user? How will you control it?"

- “Who is your first target segment and how will you acquire them cheaply?”
- “What is the biggest risk and what did you change because of it?”